

# Hyperparameter Tuning in Pytorch

Presented by Taehee Kim  
21.02.18

# Hyperparameters

- Model-free hyperparameters
  - Learning rate
  - Batch size per gpu
  - Training epoch
  - Learning rate scheduler(Warm up steps, lambda, step size ...)
  - Optimizer (beta1, beta2 in Adam)
  - Weight initialization
  - Early stop strategy
  - Regularization
  - Dropout
  - Perturbation or noise for an input
  - ...

# Hyperparameters

- Model hyperparameters
  - Kernel size
  - number of layer
  - number of hidden units
  - number of embedding units
  - pooling
  - activation function

# Hyperparameters

- Model-free & Model hyperparameters
  - Learning rate x Batch size per gpu x Training epoch x Learning rate scheduler (Warmup steps, lambda, step size...) x Optimizer (+beta1, beta2 in Adam) x Weight initialization x Early stop strategy x Regularization x Dropout x Kernel size x number of layer x number of hidden units x number of embedding units x pooling layer x activation function x....
  - [training time for a model](#)

# Hyperparameters


- ???
  - Accumulation steps
  - Random seed
  - Number of evaluation

# Hyperparameters

- Accumulation steps
  - 16 batch with no accumulation vs 4 batch with 4 accumulation step, which can result in better performance?


# Hyperparameters

- Random Seed
  - There is a promising seeds
  - These seeds can be distinguished early in training



# Hyperparameters

- Number of evaluation
  - Expected validation performance as the number of evaluation increases





# How to control randomness?

- `random.seed()`
- `np.random.seed()`
- `torch.manual_seed()`
- `torch.cuda.manual_seed()` / `torch.cuda.manual_seed_all()`
- `torch.backends.cudnn.deterministic = True`
- `torch.backends.cudnn.benchmark = False`
- [torch.set\\_deterministic\(\)](#)
- If you use CUDA tensors, we need to set the environment variable `CUBLAS_WORKSPACE_CONFIG` according to [CUDA documentation](#)

Note: The non-deterministic behavior of multi-stream execution is due to library optimizations in selecting internal workspace for the routines running in parallel streams. To avoid this effect user can either:

- provide a separate workspace for each used stream using the `cublasSetWorkspace()` function, or
- have one cuBLAS handle per stream, or
- use `cublasLtMatmul()` instead of `*gemm*()` family of functions and provide user owned workspace, or
- set a debug environment variable `CUBLAS_WORKSPACE_CONFIG` to `":16:8"` (may limit overall performance) or `":4096:8"` (will increase library footprint in GPU memory by approximately 24MiB).

# Hyperparameter Tuning

- [Hyperparameters Optimization](#)
  - Grid Search
  - Random Search
  - Bayesian
- Priority
  - For me, 1 tier = Learning rate / 1.5 tier: Batch size
  - [Hyperparameters importance are \(as for Andrew Ng\):](#) Learning rate, mini-batch size, momentum beta, number of hidden units, number of layers, learning rate decay, regularization lambda, activation functions, beta1 & beta2 in Adam

# Hyperparameter Tuning

- [Hyperparameters Optimization](#)
  - Grid Search
  - Random Search
  - Bayesian
- Priority
  - For me, 1 tier = Learning rate / 1.5 tier: Batch size
  - [Hyperparameters importance are \(as for Andrew Ng\):](#) Learning rate, mini-batch size, momentum beta, number of hidden units, number of layers, learning rate decay, regularization lambda, activation functions, beta1 & beta2 in Adam

# Hyperparameter Tuning


- (Maybe for researchers) Best practice for hyperparameters optimization
  - [Learning rate, Learning rate, Learning rate](#)
  - Step 1 (coarse-grained)
    - Turn off learning rate scheduler and train a model
    - Find a learning rate that makes loss low at the early stage of training compared to other learning rates

# Hyperparameter Tuning

- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 2 (fine-grained)
    - Turn on learning rate scheduler and train a model
    - Find a learning rate that makes loss low at the early stage of training compared to other learning rates
    - Random search around the learning rate found in step 1
    - Make possible batch sizes larger


# Hyperparameter Tuning

- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 3 (fine-grained)
    - Look at the training learning curves




# Hyperparameter Tuning

- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 3 (fine-grained)
    - Look at the training learning curves



# Hyperparameter Tuning


- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 4 (fine-grained)
    - Look at the training & validation learning curves






# Hyperparameter Tuning

- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 4 (fine-grained)
    - Look at the training & validation learning curves



# Hyperparameter Tuning

- (Maybe for researchers) Best practice for hyperparameters optimization
  - Step 4 (fine-grained)
    - Look at the training & validation learning curves



# Conclusion

- Tips for Geeks who struggle for state-of-the-art performance or want to beat competitors
  - Train your model with the largest batch size that memory allows in single gpu
  - Evaluate as many checkpoints as possible
  - Before the test, combine train set and validation set and train the model with the combined dataset
  - Do ensemble
  - The hyperparameters configuration of competitors who use similar model is a good starting point
  - Use MLOps tool (e.g., [wandb](#))